

Code-Factory System Architecture

Version: 1.0.0

Status: Draft

Last Updated: 2026-01-07

Owner: Code-Factory Core Team

Table of Contents

1. Executive Summary
2. Architecture Principles
3. System Overview
4. Component Architecture
5. Data Flow
6. Storage Model
7. Integration Points
8. Security Architecture
9. Deployment Model
10. Scalability & Performance

Executive Summary

Code-Factory is a **spec-driven software development system** that transforms natural language specifications into production-quality code through AI-powered workflows. The system is designed as a **single, zero-dependency Go binary** with an elegant terminal UI, supporting both local (Ollama) and cloud-based (OpenAI, Claude) LLM providers.

Key Characteristics

- **Distribution:** Single binary, no dependencies
- **Interface:** Beautiful TUI (Charm.sh/Bubble Tea) + optional web mirror
- **Storage:** Git-native flat files (markdown + reports)
- **AI:** Pluggable LLM backend (Ollama embedded or BYOK)
- **Integration:** Optional GitHub OAuth for seamless workflow
- **Target:** Individual developers and small teams

Design Goals

1. **Effortless:** Install and productive in 60 seconds
2. **Beautiful:** Terminal-first with aesthetic TUI
3. **Intelligent:** AI-powered but human-in-the-loop
4. **Collaborative:** Git-native, PR-ready outputs
5. **Flexible:** Works offline, online, with any LLM

Architecture Principles

1. Simplicity Over Complexity

Rationale: Developers want tools that “just work” without extensive setup.

Implementation:

- Single binary distribution (no package managers, no dependencies)
- Sensible defaults for everything
- Auto-detection and auto-configuration where possible
- Progressive disclosure of advanced features

2. Git-Native Storage

Rationale: Developers already know git; don’t introduce new data models.

Implementation:

- All specifications are markdown files in `/contracts/`
- All reports are markdown files in `/reports/`
- No databases, no proprietary formats
- Versionable, diffable, reviewable via standard git tools

3. Terminal-First Experience

Rationale: Developers live in terminals; GUI apps break flow.

Implementation:

- Primary interface is Charm.sh TUI (canvas-style, not line-by-line)
- Rich interactions: forms, menus, progress indicators
- Optional web UI as mirror for stakeholders
- Full keyboard navigation, no mouse required

4. Privacy & Security by Default

Rationale: Code is sensitive; developers value privacy.

Implementation:

- Prefer local Ollama over cloud LLMs
- Secrets in OS keyring, never in config files
- Explicit consent for any external communication
- Audit log of all AI interactions

5. Human-in-the-Loop

Rationale: AI is powerful but not infallible; humans must validate.

Implementation:

- All AI outputs require human review
- Clear diff views before any code changes
- Undo/rollback mechanisms
- Rescue mode for when AI goes off-rails

6. Extensibility Through Simplicity

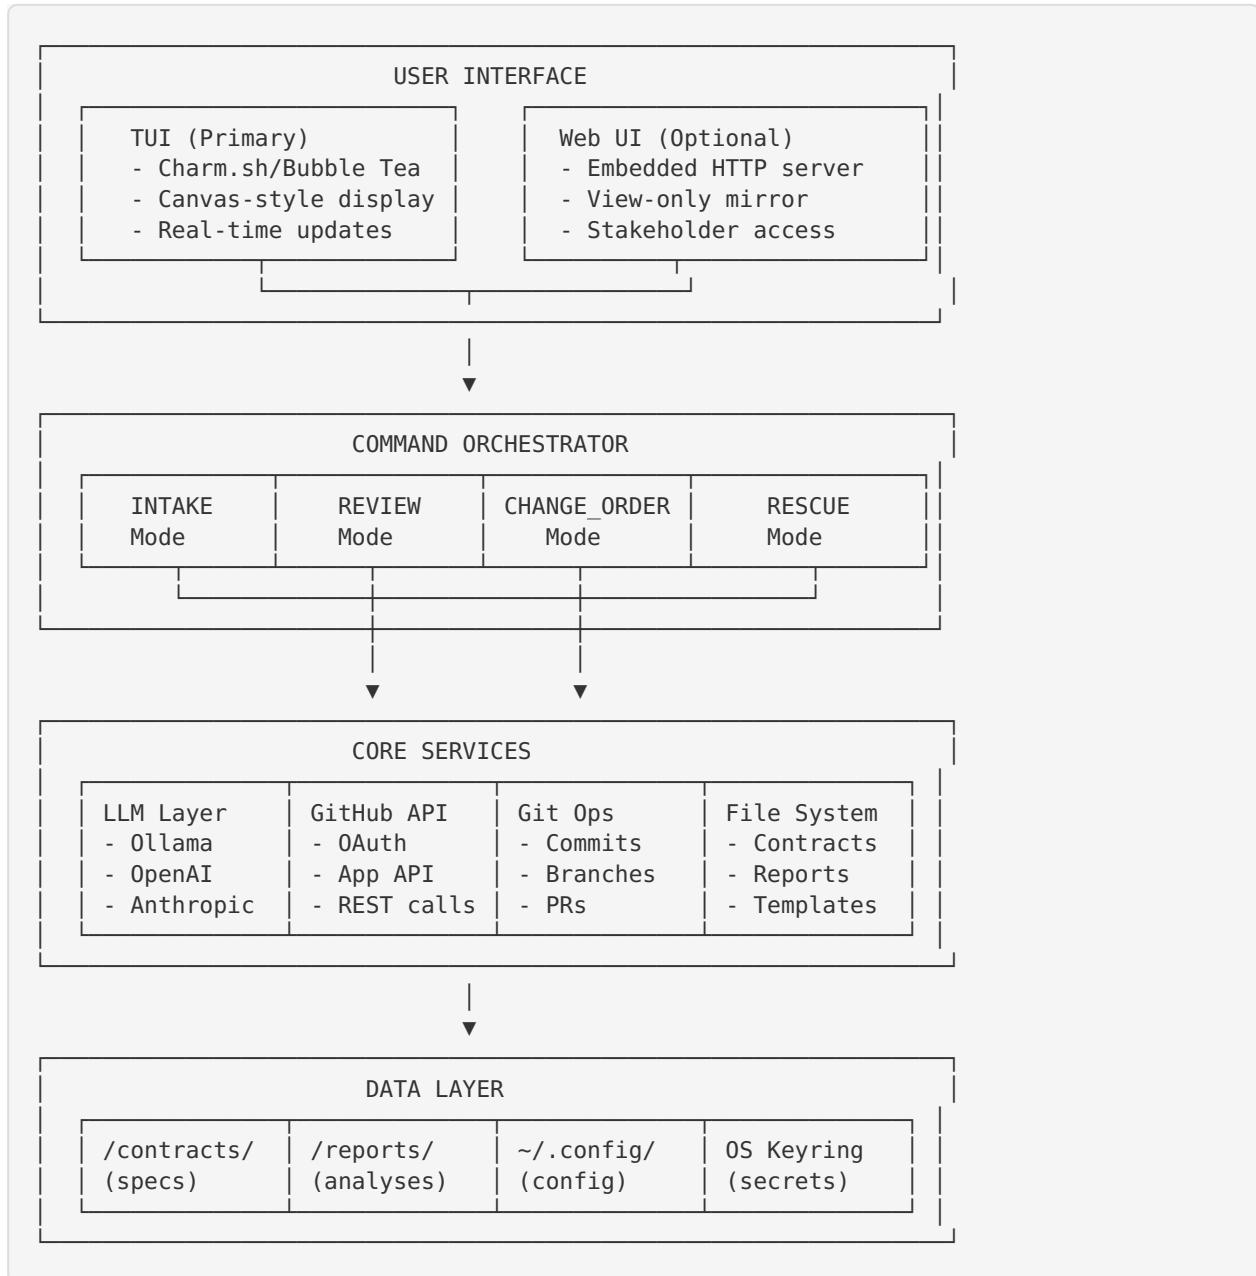
Rationale: Advanced users want to customize without fighting the tool.

Implementation:

- Plain text files enable any editor/tool integration
- Plugin system via executable hooks (future)
- Configuration in YAML, scriptable with any language
- REST API for web UI enables custom clients

System Overview

High-Level Architecture



Component Layers

Layer 1: User Interface

- **TUI:** Primary interface for all operations
- **Web UI:** Optional read-only mirror for non-technical stakeholders

Layer 2: Command Orchestrator

- **Mode Handlers:** INTAKE, REVIEW, CHANGE_ORDER, RESCUE
- **Workflow Engine:** State machine driving each mode
- **Validation:** Pre/post checks for operations

Layer 3: Core Services

- **LLM Layer:** Abstraction over multiple LLM providers
- **GitHub API:** OAuth, App API, REST operations
- **Git Ops:** Local git operations via go-git
- **File System:** Safe file I/O with rollback

Layer 4: Data Layer

- **Contracts:** Markdown specs (git-tracked)
- **Reports:** Analysis outputs (git-tracked)
- **Config:** User preferences (YAML)
- **Secrets:** OS keyring integration

Component Architecture

1. TUI Component

Technology: Charm.sh Bubble Tea + Lipgloss + Bubbles

Architecture:

```
internal/tui/
├── app.go          # Main TUI application
└── models/
    ├── base.go      # Base model with common state
    ├── intake.go     # INTAKE mode model
    ├── review.go     # REVIEW mode model
    ├── change_order.go # CHANGE_ORDER mode model
    └── rescue.go     # RESCUE mode model
└── components/
    ├── header.go    # Top status bar
    ├── footer.go     # Help/shortcuts bar
    ├── sidebar.go    # Mode selector
    ├── editor.go      # Spec editor widget
    ├── diff_viewer.go # Code diff display
    ├── file_tree.go   # File browser
    └── progress.go    # Progress indicators
└── styles/
    └── theme.go      # Color schemes
└── utils/
    └── layout.go     # Layout calculations
```

Key Features:

- **Canvas-style rendering:** Full screen, not line-by-line
- **Real-time updates:** Progress bars, spinners, status updates
- **Keyboard-driven:** No mouse required (but supported)
- **Responsive layout:** Adapts to terminal size
- **Syntax highlighting:** Via chroma library

Example Model Structure:

```

type IntakeModel struct {
    // State
    step      IntakeStep      // Current step in wizard
    spec      *Specification   // Spec being created
    llmResponse string       // LLM output
    err       error          // Error state

    // UI components
    textInput  textinput.Model // Text input widget
    textArea   textarea.Model // Multi-line editor
    progress   progress.Model // Progress bar
    viewport   viewport.Model // Scrollable area

    // Services
    llm        LLMService
    git        GitService
}

func (m IntakeModel) Init() tea.Cmd {
    return textinput.Blink
}

func (m IntakeModel) Update(msg tea.Msg) (tea.Model, tea.Cmd) {
    switch msg := msg.(type) {
    case tea.KeyMsg:
        return m.handleKeypress(msg)
    case llmResponseMsg:
        return m.handleLLMResponse(msg)
    case tea.WindowSizeMsg:
        return m.handleResize(msg)
    }
    return m, nil
}

func (m IntakeModel) View() string {
    return lipgloss.JoinVertical(
        lipgloss.Top,
        m.renderHeader(),
        m.renderContent(),
        m.renderFooter(),
    )
}

```

2. LLM Layer

Design: Provider-agnostic interface with pluggable backends

Architecture:

```

internal/llm/
└── interface.go      # LLM interface definition
  └── ollama/
    ├── client.go      # Ollama API client
    └── streaming.go   # Streaming support
  └── openai/
    ├── client.go      # OpenAI API client
    └── functions.go   # Function calling
  └── anthropic/
    ├── client.go      # Claude API client
  └── google/
    ├── client.go      # Gemini API client
  └── azure/
    ├── client.go      # Azure OpenAI
  └── custom/
    ├── client.go      # Custom endpoint
  └── provider.go      # Provider factory
  └── prompt_builder.go # Prompt construction
  └── response_parser.go # Response parsing

```

Interface:

```

type LLMService interface {
    // Generate text completion
    Generate(ctx context.Context, req GenerateRequest) (*GenerateResponse, error)
    // Stream text completion
    GenerateStream(ctx context.Context, req GenerateRequest) (<-chan GenerateChunk, error)
    // List available models
    ListModels(ctx context.Context) ([]Model, error)
    // Get provider info
    Provider() string
}

type GenerateRequest struct {
    Prompt      string
    System      string           // System prompt
    Temperature float64
    MaxTokens   int
    Stop        []string
    Context     map[string]string // Additional context
}

type GenerateResponse struct {
    Text        string
    TokensUsed int
    FinishReason string
    Model       string
}

```

Provider Selection Logic:

```

func NewLLMService(cfg *config.LLMConfig) (LLMService, error) {
    switch cfg.Provider {
    case "ollama":
        return ollama.NewClient(cfg.Endpoint, cfg.Model)
    case "openai":
        apiKey, _ := secrets.Get(cfg.APIKeyRef)
        return openai.NewClient(apiKey, cfg.Model)
    case "anthropic":
        apiKey, _ := secrets.Get(cfg.APIKeyRef)
        return anthropic.NewClient(apiKey, cfg.Model)
    // ... other providers
    default:
        return nil, fmt.Errorf("unknown provider: %s", cfg.Provider)
    }
}

```

Prompt Engineering:

```

type PromptBuilder struct {
    mode      string
    context   *Context
    template  *template.Template
}

func (pb *PromptBuilder) Build() string {
    // Load mode-specific template
    tmpl := pb.loadTemplate(pb.mode)

    // Inject context
    data := map[string]interface{}{
        "Mode":          pb.mode,
        "UserInput":     pb.context.UserInput,
        "CodeContext":   pb.context.CodeSnippets,
        "Spec":          pb.context.ExistingSpec,
    }

    var buf bytes.Buffer
    tmpl.Execute(&buf, data)
    return buf.String()
}

```

3. GitHub Integration

Architecture:

```

internal/github/
├── oauth.go          # Device flow OAuth
├── app.go            # GitHub App API
├── client.go         # REST API client
└── operations/
    ├── branch.go       # Branch operations
    ├── commit.go        # Commit operations
    ├── pr.go            # Pull request operations
    └── issue.go         # Issue operations
└── webhook.go        # Webhook handling (future)
└── sync.go           # Bi-directional sync

```

Key Operations:

```

type GitHubService interface {
    // OAuth
    InitiateDeviceFlow() (*DeviceCode, error)
    PollForToken(deviceCode string) (string, error)

    // Repository
    ListRepos() ([]Repository, error)
    GetRepo(owner, name string) (*Repository, error)

    // Branch & Commit
    CreateBranch(repo, branch, baseBranch string) error
    Commit(repo, branch string, files []FileChange) error

    // Pull Request
    CreatePR(repo string, pr *PullRequest) (*PullRequest, error)
    UpdatePR(repo string, prNumber int, updates *PRUpdate) error
    MergePR(repo string, prNumber int) error

    // Issues
    CreateIssue(repo string, issue *Issue) (*Issue, error)
}

```

PR Creation Workflow:

```

func (s *GitHubService) CreateChangeOrderPR(
    repo string,
    spec *Specification,
    changes []CodeChange,
) (*PullRequest, error) {
    // 1. Create feature branch
    branchName := fmt.Sprintf("factory/change-%s", spec.ID)
    if err := s.CreateBranch(repo, branchName, "main"); err != nil {
        return nil, err
    }

    // 2. Commit changes
    files := make([]FileChange, len(changes))
    for i, change := range changes {
        files[i] = FileChange{
            Path:    change.Path,
            Content: change.NewContent,
        }
    }
    if err := s.Commit(repo, branchName, files); err != nil {
        return nil, err
    }

    // 3. Create PR
    pr := &PullRequest{
        Title:  spec.Title,
        Body:   s.generatePRBody(spec, changes),
        Head:   branchName,
        Base:   "main",
        Labels: []string{"code-factory", "automated"},
    }
    return s.CreatePR(repo, pr)
}

```

4. Mode Handlers

Each mode is a self-contained workflow with its own state machine.

Architecture:

```
internal/modes/
└── intake/
    ├── handler.go      # Main handler
    ├── wizard.go       # Step-by-step wizard
    ├── spec_generator.go # LLM-powered spec gen
    └── validator.go    # Spec validation
└── review/
    ├── handler.go      # Code analysis
    ├── analyzer.go     # Report generation
    ├── reporter.go     # Improvement suggestions
    └── suggestions.go
└── change_order/
    ├── handler.go      # Change planning
    ├── planner.go       # Code generation
    ├── implementer.go   # PR creation
    └── pr_creator.go
└── rescue/
    ├── handler.go      # Issue detection
    ├── detector.go     # Problem diagnosis
    ├── debugger.go     # Solution generator
    └── fixer.go
```

Handler Interface:

```
type ModeHandler interface {
    // Initialize mode
    Init(ctx context.Context) error

    // Run the mode workflow
    Run(ctx context.Context, input ModeInput) (*ModeOutput, error)

    // Get current state
    State() ModeState

    // Handle user input
    HandleInput(input string) error
}
```

Example: INTAKE Mode Flow:

```

type IntakeHandler struct {
    llm    LLMService
    git    GitService
    github GitHubService
    state  IntakeState
}

func (h *IntakeHandler) Run(ctx context.Context, input ModeInput) (*ModeOutput, error) {
    // Step 1: Gather requirements
    requirements, err := h.gatherRequirements(ctx, input)
    if err != nil {
        return nil, err
    }

    // Step 2: Generate spec via LLM
    spec, err := h.generateSpec(ctx, requirements)
    if err != nil {
        return nil, err
    }

    // Step 3: User review & edit
    spec, err = h.reviewSpecWithUser(ctx, spec)
    if err != nil {
        return nil, err
    }

    // Step 4: Save to /contracts/
    if err := h.saveSpec(spec); err != nil {
        return nil, err
    }

    // Step 5: Commit to git
    if err := h.git.Commit("Add spec: " + spec.Title); err != nil {
        return nil, err
    }

    // Step 6: Optional GitHub push
    if input.PushToGitHub {
        if err := h.github.PushSpec(spec); err != nil {
            return nil, err
        }
    }

    return &ModeOutput{
        Success: true,
        Spec:    spec,
        Message: "Specification created successfully",
    }, nil
}

```

5. Storage & File System

Storage Strategy: Git-native flat files

Directory Structure:

```

project-root/
  contracts/
    README.md          # All specifications
    specs/
      feature-001.md   # Index & guide
      feature-002.md
      ...
    architecture/
      system-design.md # Individual features
      data-model.md    # High-level design
      api-spec.md      # Data structures
      decisions/
        adr-001-use-go.md # API contracts
        adr-002-tui.md   # Architecture decisions
    reports/
      review-2026-01-07.md # Generated reports
      analysis-dashboard.md # Review outputs
      change-log.md        # Aggregated metrics
    .factory/             # History of changes
    cache/                # Factory metadata (gitignored)
    temp/                 # LLM response cache
    .gitignore            # Temporary files

```

File Operations:

```

type FileService interface {
    // Read
    ReadSpec(path string) (*Specification, error)
    ListSpecs(dir string) ([]*Specification, error)

    // Write
    WriteSpec(spec *Specification) error
    UpdateSpec(spec *Specification) error
    DeleteSpec(path string) error

    // Reports
    WriteReport(report *Report) error

    // Templates
    LoadTemplate(name string) (*template.Template, error)
}

```

Spec Format:

```

---
id: feature-001
title: User Authentication
status: draft
created: 2026-01-07
updated: 2026-01-07
author: johndoe
tags: [auth, security, backend]
---

# User Authentication

## Overview
This specification defines the user authentication system...

## Requirements
- Users must be able to sign up with email
- Passwords must be hashed with bcrypt
- JWT tokens for session management

## API Endpoints
- POST /api/auth/signup
- POST /api/auth/login
- POST /api/auth/logout

## Security Considerations
...

```

Git Operations:

```

type GitService interface {
    // Status
    Status() (*GitStatus, error)

    // Basic operations
    Add(files []string) error
    Commit(message string) error
    Push(remote, branch string) error
    Pull(remote, branch string) error

    // Branching
    CreateBranch(name string) error
    SwitchBranch(name string) error

    // History
    Log(limit int) ([]*Commit, error)
    Diff(ref1, ref2 string) (*Diff, error)
}

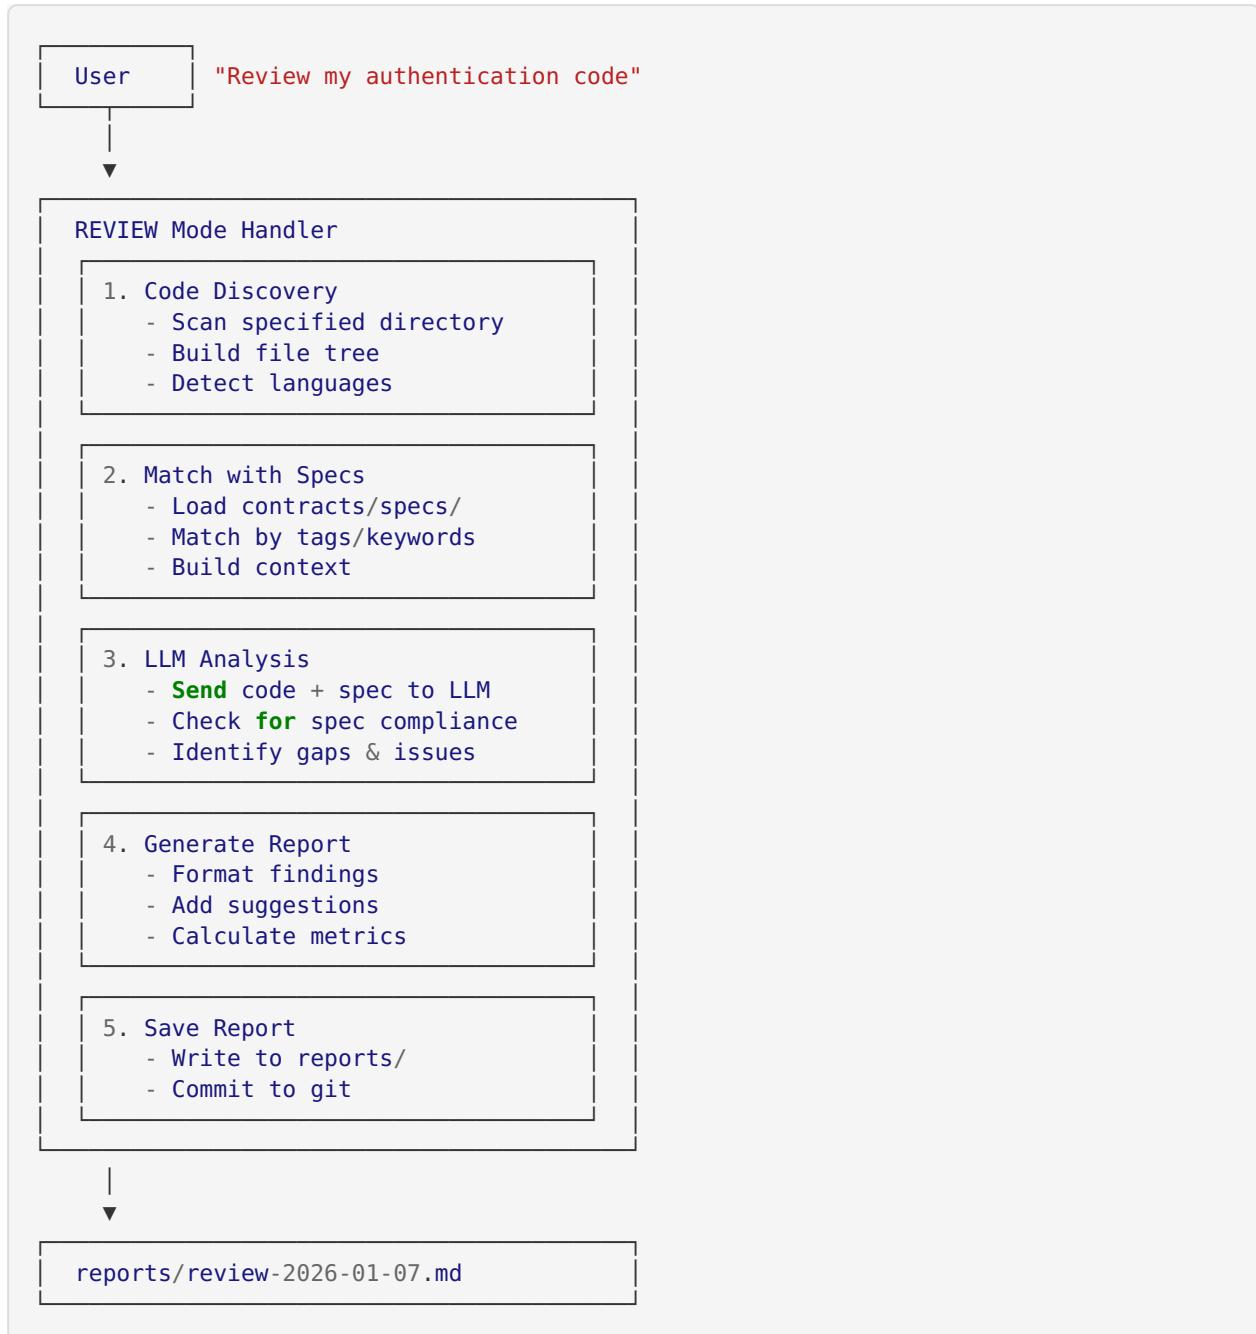
```

Data Flow

INTAKE Mode Flow



REVIEW Mode Flow



CHANGE_ORDER Mode Flow



Storage Model

Configuration Files

`~/.config/factory/config.yaml`:

```

version: "1.0.0"
user:
  name: "John Doe"
  email: "john@example.com"

llm:
  provider: "ollama"
  endpoint: "http://localhost:11434"
  model: "llama3.2:latest"
  options:
    temperature: 0.7
    num_ctx: 8192

github:
  enabled: true
  username: "johndoe"
  token_ref: "github.oauth.token"
  app_installed: true

ui:
  theme: "auto" # auto, light, dark
  animations: true
  syntax_highlight: true

projects:
  current: "/home/user/projects/my-app"
  recent:
    - "/home/user/projects/my-app"
    - "/home/user/projects/another-app"

```

Secrets (OS Keyring)

Stored securely, never in config files:

```

Service: "factory"
Accounts:
  - "llm.openai.key"      → "sk-..."
  - "llm.anthropic.key"   → "sk-ant-..."
  - "github.oauth.token"  → "ghp_..."
  - "github.installation.token" → "ghs_..."

```

Project Files

contracts/specs/template.md:

```

---
id: {auto-generated}
title: {Feature Title}
status: draft | approved | implemented | deprecated
created: {ISO date}
updated: {ISO date}
author: {username}
tags: [{tag1}, {tag2}]
priority: low | medium | high | critical
---

# {Feature Title}

## Overview
Brief description of the feature...

## Requirements

### Functional
- Requirement 1
- Requirement 2

### Non-Functional
- Performance targets
- Security requirements

## API Specification
Detailed API contracts...

## Data Model
Database schema, data structures...

## Implementation Notes
Technical considerations...

## Testing Strategy
How to verify implementation...

## Related Specs
- [Link to related spec](./other-spec.md)

```

Integration Points

External Services

1. Ollama (optional)

- Endpoint: http://localhost:11434
- Protocol: HTTP REST
- Auth: None (local)

2. OpenAI (optional)

- Endpoint: https://api.openai.com/v1
- Protocol: HTTP REST
- Auth: Bearer token

3. **Anthropic** (optional)
 - Endpoint: <https://api.anthropic.com/v1>
 - Protocol: HTTP REST
 - Auth: x-api-key header
4. **GitHub** (optional)
 - Endpoint: <https://api.github.com>
 - Protocol: HTTP REST
 - Auth: OAuth device flow

Internal Communication

TUI ↔ Core Services:

- Protocol: In-process function calls (Go)
- Pattern: Command pattern with async execution
- Error handling: Result types (value, error)

Web UI ↔ Core Services:

- Protocol: HTTP REST + WebSocket
 - API: RESTful JSON
 - Real-time updates: WebSocket for progress
-

Security Architecture

Threat Model

Assets:

- Source code
- API keys (OpenAI, GitHub)
- OAuth tokens
- Specifications (may contain sensitive info)

Threats:

1. API key leakage
2. Unauthorized access to GitHub repos
3. Malicious code injection via LLM
4. Credential theft from config files

Security Controls

1. **Secret Management**
 - OS keyring for sensitive data
 - No secrets in config files or logs
 - Encrypted fallback for systems without keyring
2. **Network Security**
 - HTTPS for all external APIs
 - Certificate validation
 - Timeout and retry limits
3. **Input Validation**
 - Sanitize all user inputs

- Validate LLM outputs before execution
- Path traversal prevention

4. Audit Logging

- Log all LLM interactions
- Log all GitHub API calls
- Never log secrets

5. Least Privilege

- GitHub OAuth: minimal scopes
- File system: only project directory
- Network: only necessary endpoints

Deployment Model

Distribution

Single Binary:

```
# macOS (Intel)
factory-darwin-amd64

# macOS (Apple Silicon)
factory-darwin-arm64

# Linux (x86_64)
factory-linux-amd64

# Linux (ARM64)
factory-linux-arm64

# Windows
factory-windows-amd64.exe
```

Installation:

```
# Via curl
curl -sSL https://factory.dev/install.sh | bash

# Via brew (macOS)
brew install code-factory

# Via apt (Ubuntu/Debian)
sudo apt install code-factory

# Manual
wget https://factory.dev/releases/v1.0.0/factory-linux-amd64
chmod +x factory-linux-amd64
sudo mv factory-linux-amd64 /usr/local/bin/factory
```

Configuration

Locations (in order of precedence):

1. Command-line flags: `--config=/path/to/config.yaml`

2. Environment variable: FACTORY_CONFIG
3. Project directory: ./factory/config.yaml
4. User config: ~/.config/factory/config.yaml
5. System config: /etc/factory/config.yaml

Updates

Auto-update mechanism:

```
func CheckForUpdates() (*UpdateInfo, error) {
    resp, err := http.Get("https://api.github.com/repos/ssdajoker/Code-Factory/releases/latest")
    // Parse response, compare versions
    // If newer version: prompt user to update
}

func UpdateSelf(version string) error {
    // Download new binary
    // Verify checksum
    // Replace current binary
    // Re-exec
}
```

User prompt:

New version available: v1.1.0 (current: v1.0.0)

Release notes:

- Added support **for** Azure OpenAI
- Improved spec parsing
- Bug fixes

[Update now] [Skip this version] [Remind me later]

Scalability & Performance

Performance Targets

Metric	Target	Rationale
Startup time	< 100ms	Fast enough to feel instant
UI render	60 FPS	Smooth animations
LLM response (first token)	< 2s	Maintains flow
File operations	< 50ms	No perceived lag
Git operations	< 200ms	Acceptable for local ops

Optimization Strategies

1. Lazy Loading

- Load specs on-demand, not at startup
- Parse markdown only when needed
- Cache parsed results

2. Async Operations

- LLM calls in background
- Git operations non-blocking
- Progress feedback for long operations

3. Response Streaming

- Stream LLM tokens to UI as received
- Better perceived performance
- User can read while waiting

4. Caching

- Cache LLM responses (with TTL)
- Cache parsed specs
- Cache GitHub API responses

5. Resource Limits

- Limit concurrent LLM requests (1)
- Limit file size for analysis (10MB)
- Limit number of files in review (100)

Scalability Considerations

Single User / Small Team:

- Current architecture is sufficient
- Local git + optional GitHub sync
- No server infrastructure needed

Future: Larger Teams

- Add centralized server for collaboration
- Real-time spec updates via WebSocket
- Shared LLM backend to reduce costs
- Role-based access control

Appendix

Technology Stack

Core:

- Language: Go 1.21+
- UI: Charm.sh (Bubble Tea, Lipgloss, Bubbles)
- Git: go-git
- Config: gopkg.in/yaml.v3

LLM Clients:

- Ollama: Direct HTTP client

- OpenAI: github.com/sashabaranov/go-openai
- Anthropic: Custom client
- Google: google.golang.org/genai

GitHub:

- github.com/google/go-github/v57
- golang.org/x/oauth2

Secrets:

- github.com/zalando/go-keyring

CLI:

- github.com/spf13/cobra
- github.com/spf13/viper

Future Enhancements

1. Plugin System

- Executable hooks for custom workflows
- Language-specific analyzers
- Custom LLM providers

2. Team Features

- Multi-user collaboration
- Real-time co-editing of specs
- Approval workflows

3. CI/CD Integration

- GitHub Actions for automated reviews
- Spec validation on PR
- Automatic report generation

4. Advanced AI

- Function calling for tool use
- Multi-turn conversations
- Context learning from project history

5. IDE Integration

- VSCode extension
- JetBrains plugin
- Language server protocol

Revision History

Version	Date	Changes	Author
1.0.0	2026-01-07	Initial architecture	Code-Factory Team